



✓ Space X Falcon 9 First Stage Landing Prediction

✓ Hands on Lab: Complete the Machine Learning Prediction lab

Estimated time needed: **60** minutes

Space X advertises Falcon 9 rocket launches on its website with a cost of 62 million dollars; other providers cost upward of 165 million dollars each, much of the savings is because Space X can reuse the first stage. Therefore if we can determine if the first stage will land, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against space X for a rocket launch. In this lab, you will create a machine learning pipeline to predict if the first stage will land given the data from the preceding labs.



Several examples of an unsuccessful landing are shown here:



Most unsuccessful landings are planned. Space X; performs a controlled landing in the oceans.

✓ Objectives

Perform exploratory Data Analysis and determine Training Labels

- create a column for the class
- Standardize the data
- Split into training data and test data

-Find best Hyperparameter for SVM, Classification Trees and Logistic Regression

- Find the method performs best using test data

✓ Import Libraries and Define Auxiliary Functions

```

1 !pip install numpy
2 !pip install pandas
3 !pip install seaborn
4 !pip install scikit-learn

```

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)

Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)

Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.0.2)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas) (1.5)

Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)

Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.11/dist-packages (from seaborn) (2.0.2)

Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.11/dist-packages (from seaborn) (2.2.2)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.11/dist-packages (from seaborn) (3.10.0)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.1.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.56.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.7)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (24.2)

Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.0.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.0)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.5)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (1.6.1)

Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (2.0.2)

Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.14.1)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)

We will import the following libraries for the lab

```

1 # Pandas is a software library written for the Python programming language for data manipulation and analysis.
2 import pandas as pd
3 # NumPy is a library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on these arrays.
4 import numpy as np
5 # Matplotlib is a plotting library for python and pyplot gives us a MatLab like plotting framework. We will use this in our plot.
6 import matplotlib.pyplot as plt
7 #Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical plots.
8 import seaborn as sns
9 # Preprocessing allows us to standardize our data
10 from sklearn import preprocessing
11 # Allows us to split our data into training and testing data
12 from sklearn.model_selection import train_test_split
13 # Allows us to test parameters of classification algorithms and find the best one
14 from sklearn.model_selection import GridSearchCV
15 # Logistic Regression classification algorithm
16 from sklearn.linear_model import LogisticRegression
17 # Support Vector Machine classification algorithm
18 from sklearn.svm import SVC
19 # Decision Tree classification algorithm
20 from sklearn.tree import DecisionTreeClassifier
21 # K Nearest Neighbors classification algorithm
22 from sklearn.neighbors import KNeighborsClassifier

```

This function is to plot the confusion matrix.

```

1 def plot_confusion_matrix(y,y_predict):
2     "this function plots the confusion matrix"
3     from sklearn.metrics import confusion_matrix
4
5     cm = confusion_matrix(y, y_predict)
6     ax= plt.subplot()
7     sns.heatmap(cm, annot=True, ax = ax); #annot=True to annotate cells
8     ax.set_xlabel('Predicted labels')
9     ax.set_ylabel('True labels')
10    ax.set_title('Confusion Matrix');
11    ax.xaxis.set_ticklabels(['did not land', 'land']); ax.yaxis.set_ticklabels(['did not land', 'landed'])
12    plt.show()

```

✓ Load the dataframe

Load the data

```
1 data = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/datasets")
2
```

```
1 data.head()
```

	FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block
0	1	2010-06-04	Falcon 9	6104.959412	LEO	CCAFS SLC 40	None None	1	False	False	False	NaN	
1	2	2012-05-22	Falcon 9	525.000000	LEO	CCAFS SLC 40	None None	1	False	False	False	NaN	
2	3	2013-03-01	Falcon 9	677.000000	ISS	CCAFS SLC 40	None None	1	False	False	False	NaN	

```
1
2 X = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/datasets')
```

```
1 X.head(100)
```

	FlightNumber	PayloadMass	Flights	Block	ReusedCount	Orbit_ES-L1	Orbit_GEO	Orbit_GTO	Orbit_HEO	Orbit_ISS	...	Serial_
0	1.0	6104.959412	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
1	2.0	525.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
2	3.0	677.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	...	
3	4.0	500.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
4	5.0	3170.000000	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	...	
...	
85	86.0	15400.000000	2.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	
86	87.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	
87	88.0	15400.000000	6.0	5.0	5.0	0.0	0.0	0.0	0.0	0.0	...	
88	89.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	
89	90.0	3681.000000	1.0	5.0	0.0	0.0	0.0	0.0	0.0	0.0	...	

✓ TASK 1

Create a NumPy array from the column `Class` in `data`, by applying the method `to_numpy()` then assign it to the variable `Y`, make sure the output is a Pandas series (only one bracket `df['name of column']`).

```
1 Y = data['Class'].to_numpy()
```

✓ TASK 2

Standardize the data in `X` then reassign it to the variable `X` using the transform provided below.

```
1 transform = preprocessing.StandardScaler()
2 X = transform.fit_transform(X)
```

We split the data into training and testing data using the function `train_test_split`. The training data is divided into validation data, a second set used for training data; then the models are trained and hyperparameters are selected using the function `GridSearchCV`.

✓ TASK 3

Use the function `train_test_split` to split the data `X` and `Y` into training and test data. Set the parameter `test_size` to 0.2 and `random_state` to 2. The training data and test data should be assigned to the following labels.

`X_train, X_test, Y_train, Y_test`

```
1 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

we can see we only have 18 test samples.

```
1 Y_test.shape
```

```
➦ (18,)
```

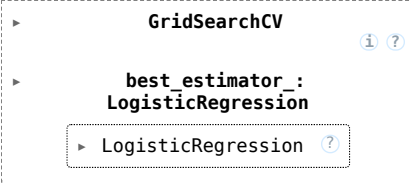
✓ TASK 4

Create a logistic regression object then create a `GridSearchCV` object `logreg_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
1 parameters = {'C':[0.01,0.1,1],
2               'penalty':['l2'],
3               'solver':['lbfgs']}
```

```
1 parameters = {'C':[0.01,0.1,1], 'penalty':['l2'], 'solver':['lbfgs']}# l1 lasso l2 ridge
2 lr = LogisticRegression()
3 logreg_cv = GridSearchCV(lr, parameters, cv=10)
4 logreg_cv.fit(X_train, Y_train)
5
```

```
➦
```



We output the `GridSearchCV` object for logistic regression. We display the best parameters using the data attribute `best_params_` and the accuracy on the validation data using the data attribute `best_score_`.

```
1 print("tuned hyperparameters :(best parameters) ",logreg_cv.best_params_)
2 print("accuracy :",logreg_cv.best_score_)
```

```
➦ tuned hyperparameters :(best parameters) {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
  accuracy : 0.8464285714285713
```

✓ TASK 5


Calculate the accuracy on the test data using the method `score`:

```
1 logreg_cv.score(X_test, Y_test)
```

```
➦ 0.8333333333333334
```

Lets look at the confusion matrix:

```
1 yhat=logreg_cv.predict(X_test)
2 plot_confusion_matrix(Y_test,yhat)
```



Examining the confusion matrix, we see that logistic regression can distinguish between the different classes. We see that the problem is false positives.

Overview:

True Postive - 12 (True label is landed, Predicted label is also landed)


False Postive - 3 (True label is not landed, Predicted label is landed)

✓ TASK 6

Create a support vector machine object then create a GridSearchCV object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary parameters.

```
1 parameters = {'kernel':('linear', 'rbf','poly','rbf', 'sigmoid'),
2               'C': np.logspace(-3, 3, 5),
3               'gamma':np.logspace(-3, 3, 5)}
4 svm = SVC()
```

1 Comienza a programar o [generar](#) con IA.




```
1 print("tuned hpyerparameters :(best parameters) ",svm_cv.best_params_)
2 print("accuracy :",svm_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'C': np.float64(1.0), 'gamma': np.float64(0.03162277660168379), 'kernel': 'sigmoi
accuracy : 0.8482142857142856
```

✓ TASK 7

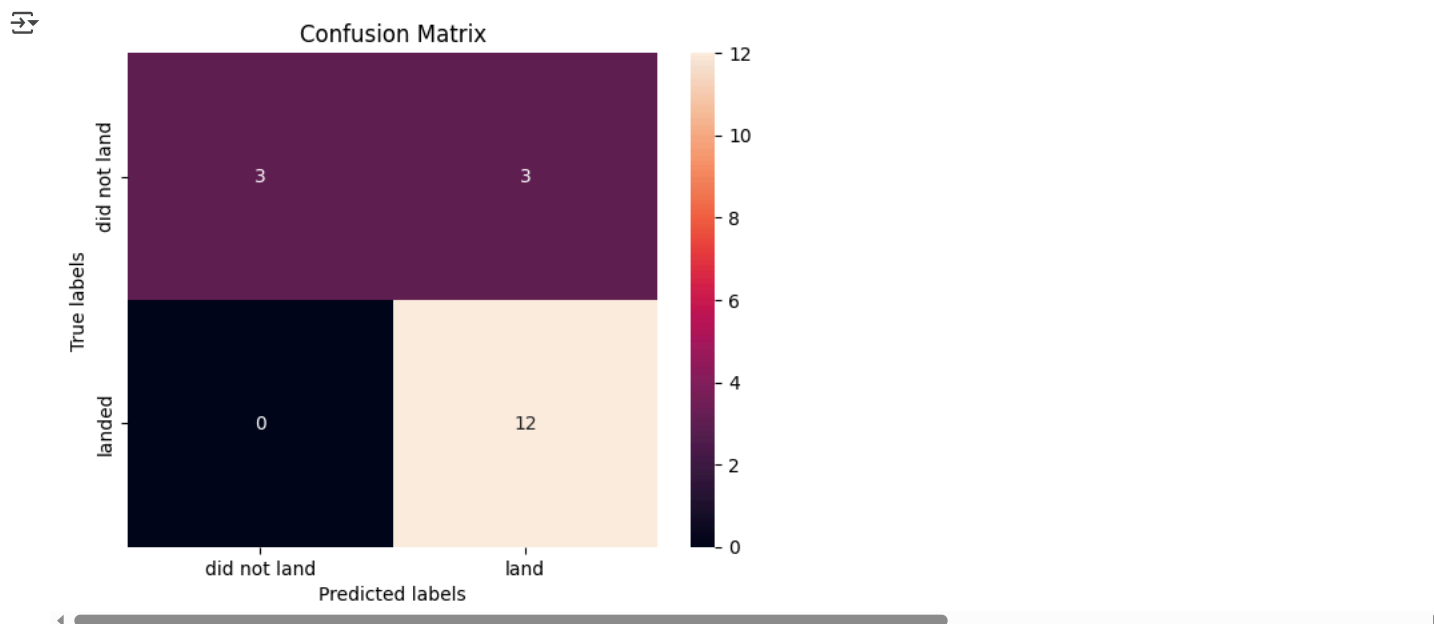
Calculate the accuracy on the test data using the method `score`:

```
1 svm_cv = GridSearchCV(svm, parameters, cv=10)
2 svm_cv.fit(X_train, Y_train)
```



We can plot the confusion matrix

```
1 yhat=svm_cv.predict(X_test)
2 plot_confusion_matrix(Y_test,yhat)
```



✓ TASK 8

Create a decision tree classifier object then create a `GridSearchCV` object `tree_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```

1 parameters = {'criterion': ['gini', 'entropy'],
2               'splitter': ['best', 'random'],
3               'max_depth': [2*n for n in range(1,10)],
4               'max_features': ['auto', 'sqrt'],
5               'min_samples_leaf': [1, 2, 4],
6               'min_samples_split': [2, 5, 10]}
7
8 tree = DecisionTreeClassifier()

1 tree_cv = GridSearchCV(tree, parameters, cv=10)
2 tree_cv.fit(X_train, Y_train)
```

```

/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py:528: FitFailedWarning:
3240 fits failed out of a total of 6480.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:
-----
3240 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 1382, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.11/dist-packages/sklearn/utils/_param_validation.py", line 98, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of DecisionTreeClassifier must be an integer or None.

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_search.py:1108: UserWarning: One or more of the test scores
nan nan nan nan nan nan
nan nan nan nan nan nan
0.82142857 0.79464286 0.81964286 0.7375 0.78035714 0.73392857
0.72678571 0.81964286 0.75178571 0.81607143 0.79107143 0.73928571
0.80535714 0.73928571 0.69464286 0.73214286 0.7625 0.75357143
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.81964286 0.79107143 0.80357143 0.82678571 0.76071429 0.88928571
0.79464286 0.83035714 0.80178571 0.83035714 0.79107143 0.74642857
0.76607143 0.79107143 0.76428571 0.81964286 0.71071429 0.77678571
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.76071429 0.84642857 0.76428571 0.76607143 0.80357143 0.80357143
0.80535714 0.80892857 0.76428571 0.83392857 0.84642857 0.78928571
0.73392857 0.80535714 0.7625 0.81785714 0.76071429 0.83035714
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.74642857 0.83214286 0.81964286 0.76428571 0.81964286 0.71964286
0.79464286 0.76607143 0.84642857 0.70535714 0.76428571 0.81785714
0.73392857 0.79285714 0.71785714 0.83214286 0.81785714 0.77678571
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.77857143 0.75 0.78928571 0.81964286 0.73571429 0.80178571
0.73392857 0.79107143 0.80535714 0.79464286 0.7625 0.74821429
0.72321429 0.79107143 0.84642857 0.7625 0.81785714 0.78928571
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.83214286 0.71964286 0.70714286 0.83571429 0.73214286 0.80714286
0.80178571 0.76071429 0.73571429 0.76071429 0.74821429 0.74642857
0.71964286 0.86071429 0.81607143 0.76785714 0.83035714 0.76607143
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.73571429 0.77678571 0.81607143 0.79107143 0.7375 0.77678571
0.74821429 0.77678571 0.8625 0.75178571 0.775 0.66607143
0.77321429 0.79107143 0.71964286 0.7625 0.81964286 0.79107143
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.7875 0.8625 0.79107143 0.86071429 0.81964286 0.79464286
0.78928571 0.73214286 0.69107143 0.72321429 0.79107143 0.84821429
0.77678571 0.83214286 0.84464286 0.77857143 0.74642857 0.775
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.80714286 0.76607143 0.73571429 0.72142857 0.81785714 0.78928571
0.80357143 0.80714286 0.7625 0.77678571 0.74107143 0.79285714
0.7625 0.81785714 0.7625 0.77678571 0.80714286 0.76071429
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.75357143 0.68214286 0.80357143 0.73928571 0.76071429 0.80535714
0.77678571 0.83392857 0.72321429 0.75178571 0.81607143 0.78928571
0.78928571 0.81964286 0.71964286 0.75178571 0.75 0.72321429
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.79107143 0.80535714 0.76607143 0.83392857 0.80714286 0.81964286
0.7625 0.84821429 0.78928571 0.83392857 0.83392857 0.7375

```

```

0.8025 0.84821429 0.78928571 0.83392857 0.83392857 0.775
0.775 0.81785714 0.77678571 0.77678571 0.73571429 0.80535714
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.88928571 0.84642857 0.73392857 0.73928571 0.77678571 0.80357143
0.76428571 0.71964286 0.76428571 0.81607143 0.80178571 0.85892857
0.80357143 0.80714286 0.77678571 0.83214286 0.80357143 0.75178571
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.74821429 0.80535714 0.73571429 0.84464286 0.81964286 0.80892857
0.73571429 0.79285714 0.70714286 0.77678571 0.75 0.81964286
0.73214286 0.78928571 0.7625 0.79107143 0.76607143 0.81964286
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.775 0.78928571 0.79107143 0.80357143 0.7625 0.77857143
0.79285714 0.84642857 0.83035714 0.80535714 0.84464286 0.775
0.81785714 0.80535714 0.81964286 0.83035714 0.71785714 0.79285714
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.79107143 0.73571429 0.70714286 0.79107143 0.77678571 0.76428571
0.70535714 0.77678571 0.74642857 0.83392857 0.80357143 0.75
0.7375 0.79107143 0.80535714 0.69464286 0.84642857 0.81071429
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.7625 0.76607143 0.76428571 0.84821429 0.73214286 0.81964286
0.775 0.70892857 0.83035714 0.77857143 0.75178571 0.81964286
0.77678571 0.79107143 0.78928571 0.81785714 0.79107143 0.74107143
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.75 0.77857143 0.74821429 0.80714286 0.775 0.77857143
0.73035714 0.83035714 0.80357143 0.8625 0.83392857 0.81964286
0.77678571 0.75 0.80535714 0.80714286 0.78928571 0.79107143
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
0.73571429 0.84821429 0.77857143 0.75 0.73392857 0.80714286
0.74642857 0.81785714 0.69464286 0.81964286 0.7625 0.78928571
0.7625 0.77678571 0.72142857 0.79464286 0.80357143 0.80714286]
warnings.warn(

```

GridSearchCV

best_estimator_:
DecisionTreeClassifier

DecisionTreeClassifier


```
1 print("tuned hyperparameters :(best parameters) ",tree_cv.best_params_)
2 print("accuracy :",tree_cv.best_score_)
```

```
tuned hyperparameters :(best parameters) {'criterion': 'gini', 'max_depth': 4, 'max_features': 'sqrt', 'min_samples_leaf':
accuracy : 0.8892857142857145
```

✓ TASK 9


Calculate the accuracy of tree_cv on the test data using the method score :

```
1 tree_cv.score(X_test, Y_test)
2
```

```
0.9444444444444444
```

We can plot the confusion matrix

```
1 yhat = tree_cv.predict(X_test)
2 plot_confusion_matrix(Y_test,yhat)
```



✓ TASK 10

Create a k nearest neighbors object then create a GridSearchCV object knn_cv with cv = 10. Fit the object to find the best parameters from the dictionary parameters .

```
1 parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
2               'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
3               'p': [1,2]}
4
5 KNN = KNeighborsClassifier()
```

```
1 KNN = KNeighborsClassifier()
2 knn_cv = GridSearchCV(KNN, parameters, cv=10)
3 knn_cv.fit(X_train, Y_train)
```